

Class 11: Structural Bioinformatics pt.2

Jennifer Thai (PID: A17893762)

Table of contents

AlphaFold Data Base (AFDB)	1
Generating your own structure predictions	2
Custom analysis of resulting models in R	5
Residue conservation from alignment file	8

AlphaFold Data Base (AFDB)

The EBI maintains the largest database of AlphaFold structure prediction models at:
<https://alphafold.ebi.ac.uk/>

From last class (before Halloween), we saw that the PDB had 244,290 (Oct 2025)

The total number of protein sequences in UniProtKB is 199,579,901

Key Point: This is a tiny fraction of sequence space that has structural coverage (0.12%)

244290 / 199579901 * 100

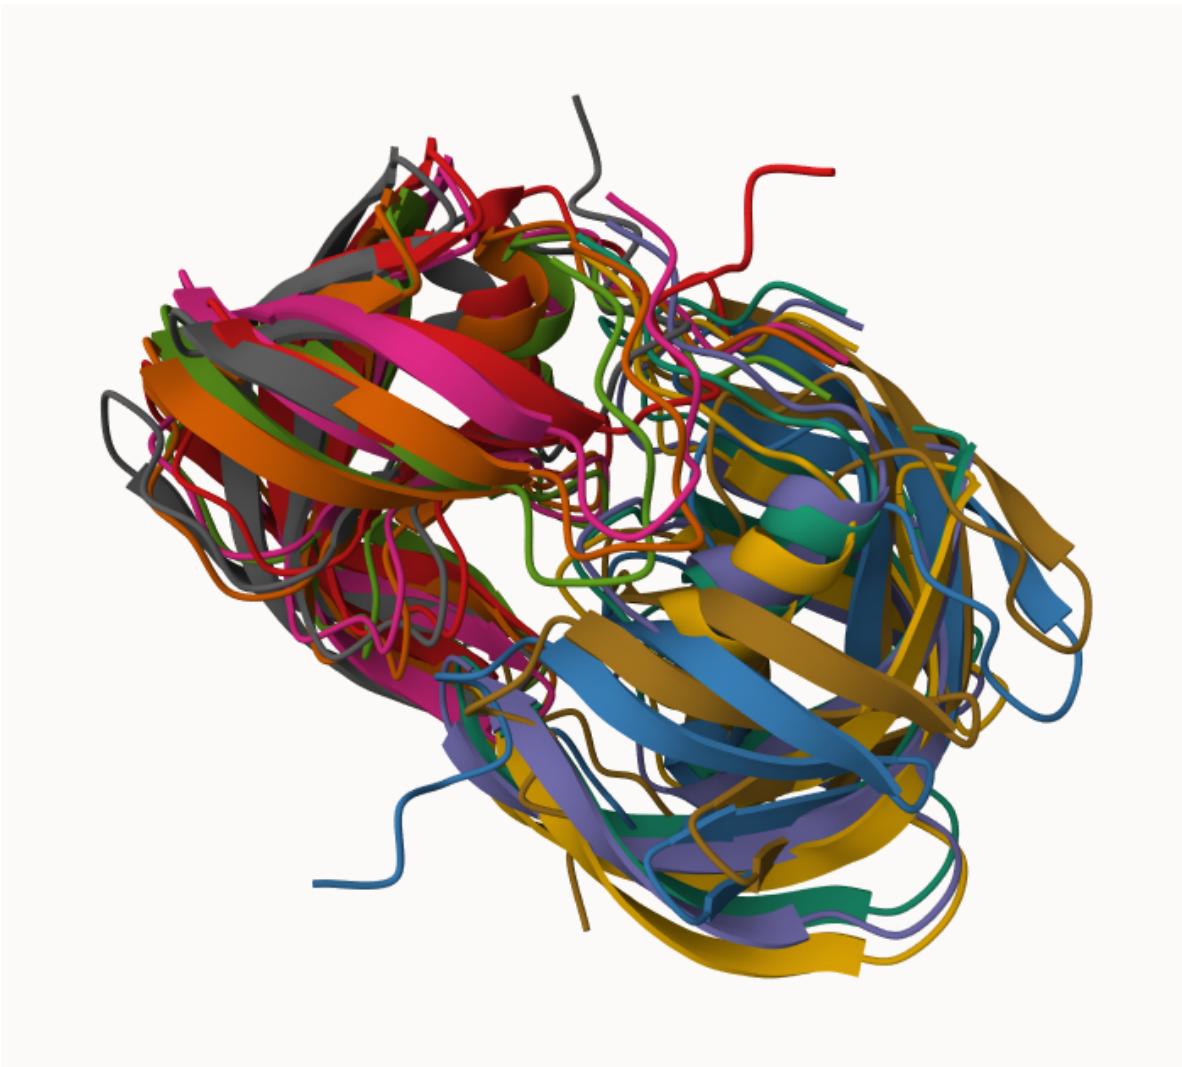
```
[1] 0.1224021
```

AFDB is attempting to address this gap...

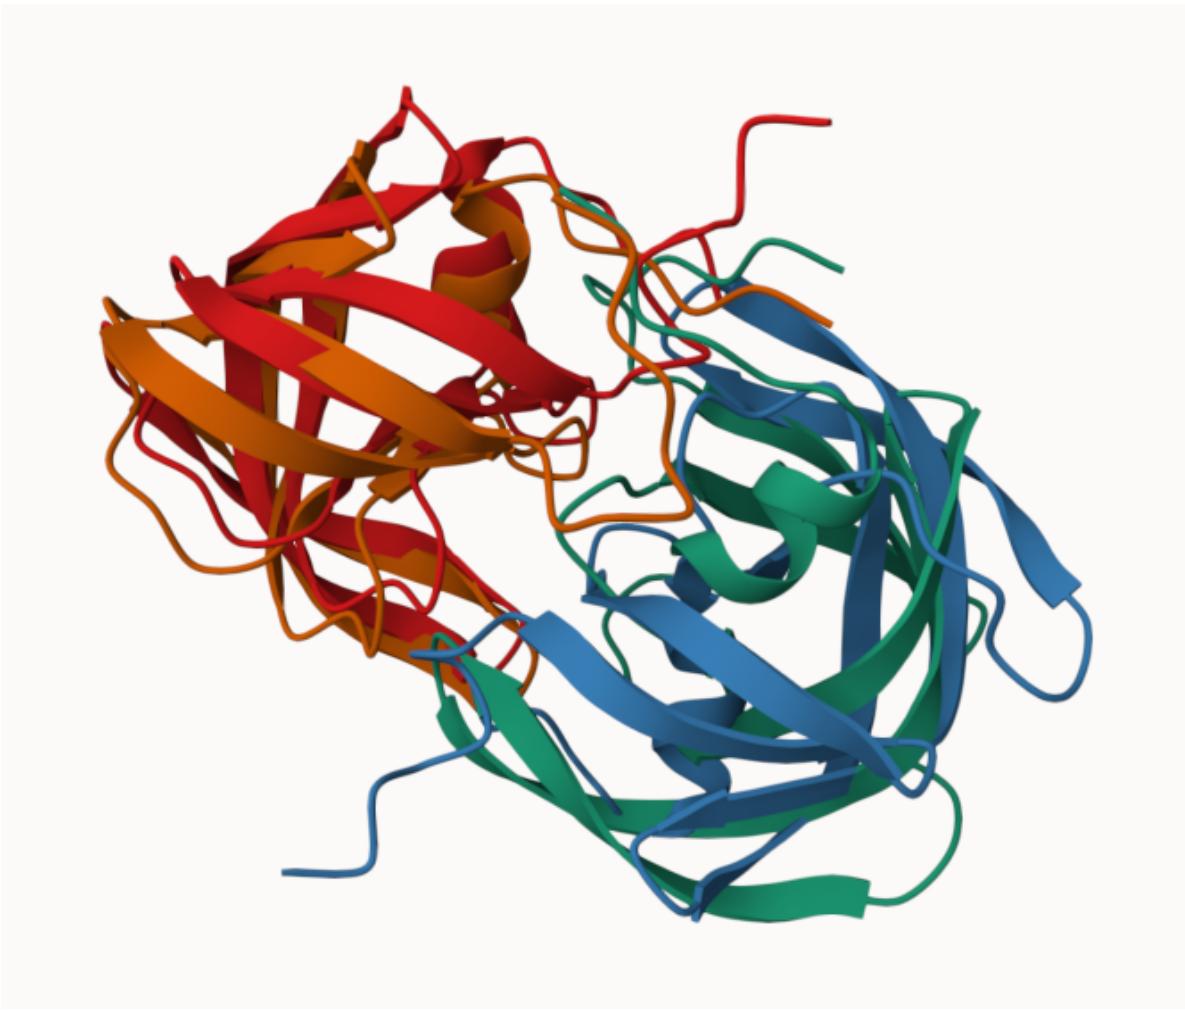
There are two “Quality Scores” form ALphaFold: one for residues (i.e. each amino acid) called **pLDDT** score. The other **PAE** score that measures the confidence in the relative position of two residues (i.e. a score for every pair of residues).

Generating your own structure predictions

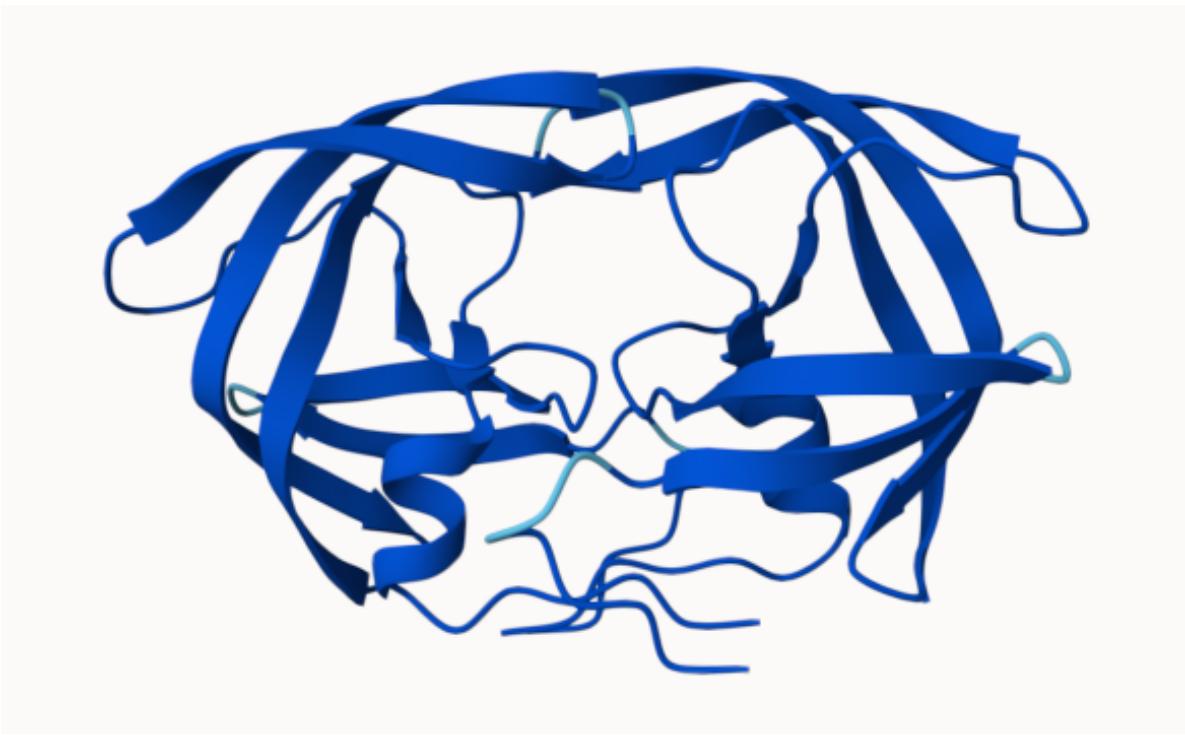
Figure of 5 generated HIV-PR models



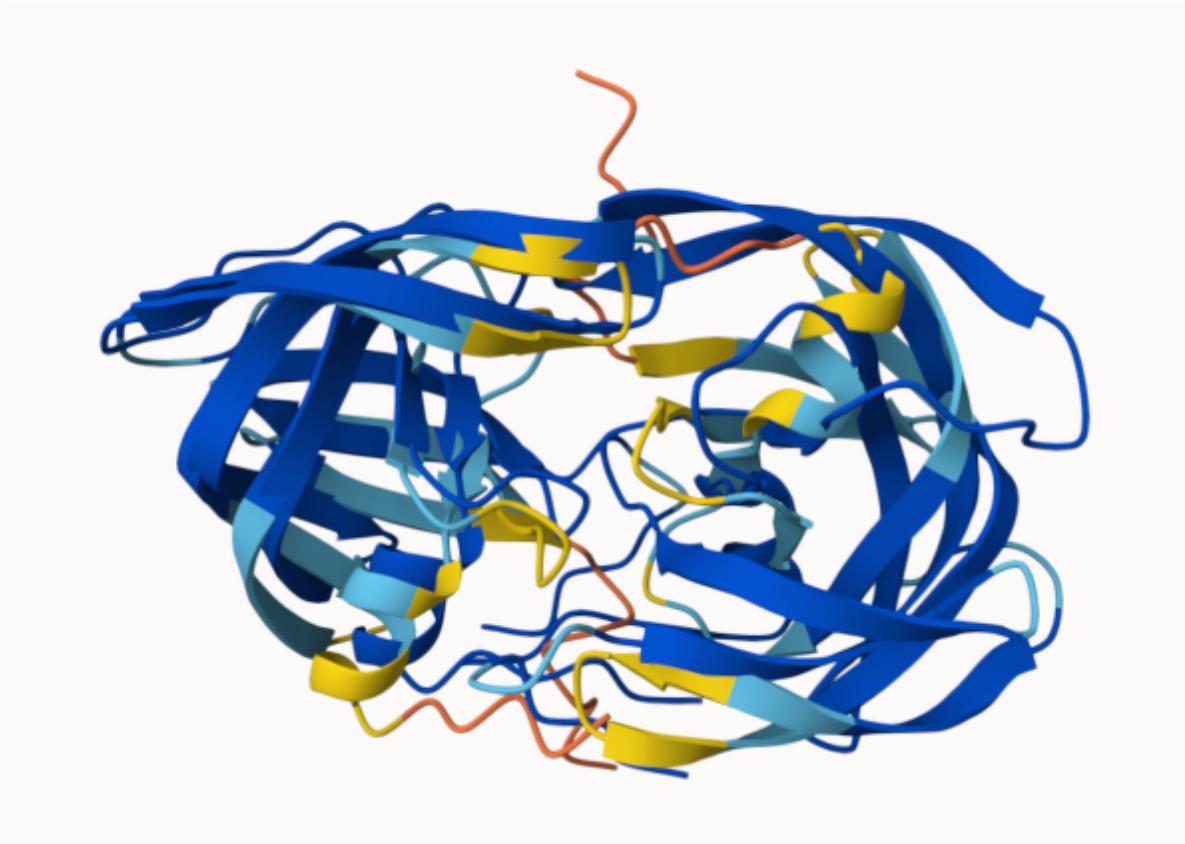
And the top ranked model colored by chain



pLDDT score for model 1



and model 5



Custom analysis of resulting models in R

Read key result files into R. The first thing I need to know is what my results directory/folder is called (i.e. its name is different for every AlphaFold run/job)

```
results_dir <- "HIVPR_dimer_23119/"
```

```
# File names for all PDB models
pdb_files <- list.files(path=results_dir,
                         pattern="*.pdb",
                         full.names = TRUE)
```

```
# Print our PDB file names
basename(pdb_files)
```

```
[1] "HIVPR_dimer_23119_unrelaxed_rank_001_alphaFold2_multimer_v3_model_4_seed_000.pdb"
```

```
[2] "HIVPR_dimer_23119_unrelaxed_rank_002_alphaFold2_multimer_v3_model_1_seed_000.pdb"
[3] "HIVPR_dimer_23119_unrelaxed_rank_003_alphaFold2_multimer_v3_model_5_seed_000.pdb"
[4] "HIVPR_dimer_23119_unrelaxed_rank_004_alphaFold2_multimer_v3_model_2_seed_000.pdb"
[5] "HIVPR_dimer_23119_unrelaxed_rank_005_alphaFold2_multimer_v3_model_3_seed_000.pdb"
```

```
library(bio3d)
```

Warning: package 'bio3d' was built under R version 4.5.2

```
m1 <- read.pdb(pdb_files[1])
m1
```

Call: read.pdb(file = pdb_files[1])

Total Models#: 1

Total Atoms#: 1514, XYZs#: 4542 Chains#: 2 (values: A B)

Protein Atoms#: 1514 (residues/Calpha atoms#: 198)

Nucleic acid Atoms#: 0 (residues/phosphate atoms#: 0)

Non-protein/nucleic Atoms#: 0 (residues: 0)

Non-protein/nucleic resid values: [none]

Protein sequence:

```
PQITLWQRPLVTIKIGGQLKEALLDTGADDTVLEEMSLPGRWPKMIGGIGGFVKVRQYD
QILIEICGHKAIGTVLVGPTPVNIIGRNLLTQIGCTLNFPQITLWQRPLVTIKIGGQLKE
ALLDTGADDTVLEEMSLPGRWPKMIGGIGGFVKVRQYDQILIEICGHKAIGTVLVGPTP
VNIIGRNLLTQIGCTLNF
```

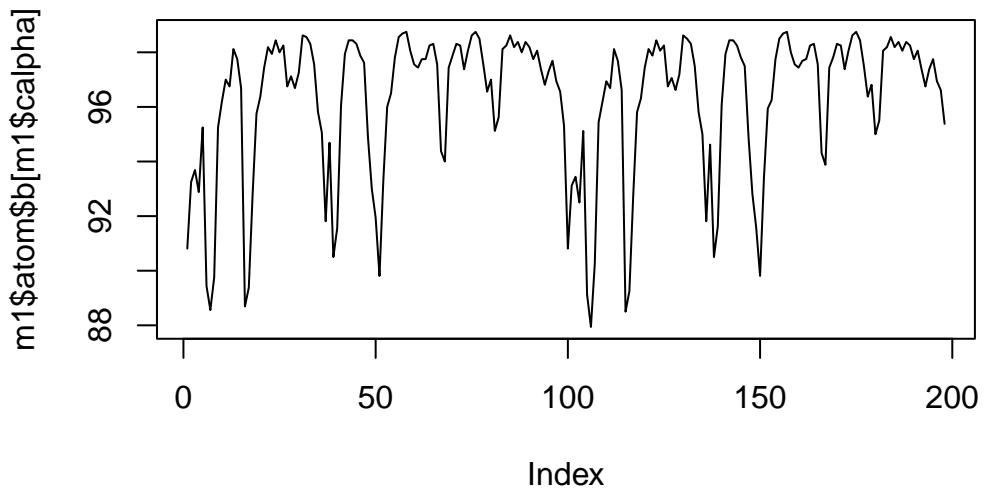
+ attr: atom, xyz, calpha, call

```
head(m1$atom)
```

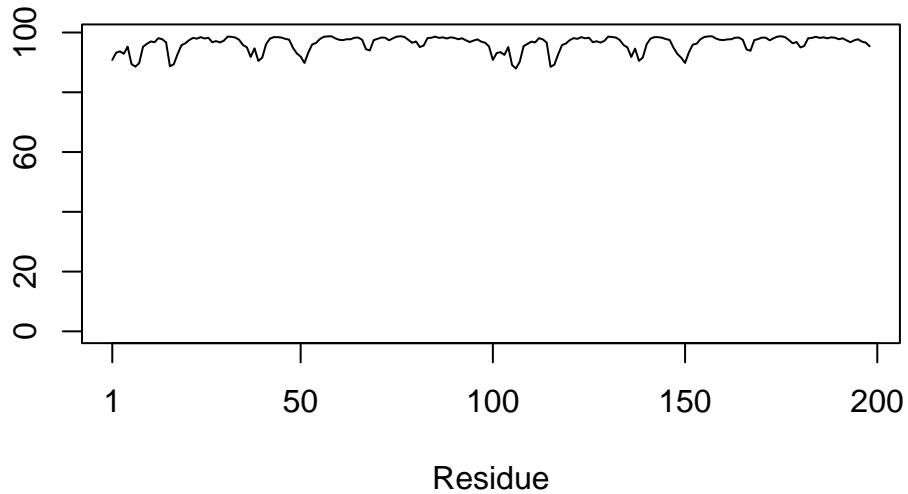
	type	eleno	elety	alt	resid	chain	resno	insert	x	y	z	o	b
1	ATOM	1	N	<NA>	PRO	A	1	<NA>	16.922	-3.898	-6.254	1	90.81
2	ATOM	2	CA	<NA>	PRO	A	1	<NA>	16.891	-2.467	-6.562	1	90.81
3	ATOM	3	C	<NA>	PRO	A	1	<NA>	16.406	-1.617	-5.395	1	90.81
4	ATOM	4	CB	<NA>	PRO	A	1	<NA>	15.930	-2.373	-7.746	1	90.81
5	ATOM	5	O	<NA>	PRO	A	1	<NA>	15.820	-2.146	-4.445	1	90.81

```
6 ATOM      6      CG <NA>    PRO      A      1      <NA> 15.031 -3.559 -7.598 1 90.81
segid elesy charge
1 <NA>      N      <NA>
2 <NA>      C      <NA>
3 <NA>      C      <NA>
4 <NA>      C      <NA>
5 <NA>      O      <NA>
6 <NA>      C      <NA>
```

```
plot(m1$atom$b[m1$calpha], typ="l")
```



```
plot.bio3d(m1$atom$b[m1$calpha], typ="l")
```



Residue conservation from alignment file

Find the large AlphaFold alignment file

```
aln_file <- list.files(path=results_dir,
                        pattern=".a3m$",
                        full.names = TRUE)
aln_file
```

```
[1] "HIVPR_dimer_23119/HIVPR_dimer_23119.a3m"
```

Read this into R

```
aln <- read.fasta(aln_file[1], to.upper = TRUE)
```

```
[1] " ** Duplicated sequence id's: 101 **"
[2] " ** Duplicated sequence id's: 101 **"
```

How many sequences are in this alignment?

```
dim(aln$ali)
```

```
[1] 5397 132
```

```
con <- consensus(aln, cutoff = 0.9)
con$seq
```

```
[1] "-"
[19] "-"
[37] "-"
[55] "-"
[73] "-"
[91] "-"
[109] "-"
[127] "-"

[1] "-"
[19] "D"
[37] "T"
[55] "G"
[73] "A"
[91] "-"
[109] "-"
[127] "-"

[1] "-"
[19] "T"
[37] "G"
[55] "A"
[73] "C"
[91] "C"
[109] "C"
[127] "C"

[1] "-"
[19] "D"
[37] "T"
[55] "G"
[73] "A"
[91] "C"
[109] "C"
[127] "C"

[1] "-"
[19] "D"
[37] "T"
[55] "G"
[73] "A"
[91] "C"
[109] "C"
[127] "C"

[1] "-"
[19] "D"
[37] "T"
[55] "G"
[73] "A"
[91] "C"
[109] "C"
[127] "C"

[1] "-"
[19] "D"
[37] "T"
[55] "G"
[73] "A"
[91] "C"
[109] "C"
[127] "C"
```

We can score residue conservation in the alignment with the conserv() function.

```
sim <- conserv(aln)
```

```
plotb3(sim[1:99], ylab="Conservation Score")
```

