

# Class 12: RNASeq analysis

Jennifer Thai (A17893762)

## Table of contents

Background . . . . .	1
Data Import . . . . .	1
Toy differential gene expression . . . . .	3
DESeq2 analysis . . . . .	8
Volcano plot . . . . .	10
Save our results . . . . .	11
Add gene annotation . . . . .	11
Pathway analysis . . . . .	13
Save our main results . . . . .	15

## Background

Today we will analyze some RNASeq data from HImes et al. on the effects of a common steroid (dexamethasone) on airway smooth muscle cells (ASM cells).

Our starting point is the “counts” data and the “metadata” that contain the count values for each genes in their different experiments (i.e. cell lines with or without the drug).

## Data Import

```
# Complete the missing code
counts <- read.csv("airway_scaledcounts.csv", row.names=1)
metadata <- read.csv("airway_metadata.csv")
```

Let's have a wee peak at these objects:

```
head(counts)
```

	SRR1039508	SRR1039509	SRR1039512	SRR1039513	SRR1039516
ENSG000000000003	723	486	904	445	1170
ENSG000000000005	0	0	0	0	0
ENSG00000000419	467	523	616	371	582
ENSG00000000457	347	258	364	237	318
ENSG00000000460	96	81	73	66	118
ENSG00000000938	0	0	1	0	2
	SRR1039517	SRR1039520	SRR1039521		
ENSG000000000003	1097	806	604		
ENSG000000000005	0	0	0		
ENSG00000000419	781	417	509		
ENSG00000000457	447	330	324		
ENSG00000000460	94	102	74		
ENSG00000000938	0	0	0		

Q1. How many genes are in this dataset?

```
nrow(counts)
```

```
[1] 38694
```

```
metadata
```

	id	dex	celltype	geo_id
1	SRR1039508	control	N61311	GSM1275862
2	SRR1039509	treated	N61311	GSM1275863
3	SRR1039512	control	N052611	GSM1275866
4	SRR1039513	treated	N052611	GSM1275867
5	SRR1039516	control	N080611	GSM1275870
6	SRR1039517	treated	N080611	GSM1275871
7	SRR1039520	control	N061011	GSM1275874
8	SRR1039521	treated	N061011	GSM1275875

Q. How many different experiments (columns in counts or rows in metadata) are there?

```
ncol(counts)
```

```
[1] 8
```

```
nrow(metadata)
```

```
[1] 8
```

Q2. How many ‘control’ cell lines do we have?

```
sum(metadata$dex == "control")
```

```
[1] 4
```

## Toy differential gene expression

To start our analysis, let’s calculate the mean counts for all genes in the “control” experiments.

1. Extract all “control” columns from the `counts` object
  2. Calculate the mean for all rows (i.e. genes) of these “control” columns
- 3-4. Do the same for “treated” 5. Compare these `control.mean` and `treated.mean` values

Q3. How would you make the above code in either approach more robust? Is there a function that could help here?

```
control inds <- metadata$dex == "control"  
control counts <- counts[ , control inds]
```

```
control means <- rowMeans(control counts)
```

Q4. Follow the same procedure for the treated samples (i.e. calculate the mean per gene across drug treated samples and assign to a labeled vector called `treated.mean`)

```
treated inds <- metadata$dex == "treated"  
treated counts <- counts[ , treated inds]
```

```
treated means <- rowMeans(treated counts)
```

Store these together for ease of bookkeeping as `meancounts`


```
meancounts <- data.frame(control means, treated means)  
head(meancounts)
```

	control.means	treated.means
ENSG000000000003	900.75	658.00
ENSG000000000005	0.00	0.00
ENSG000000000419	520.50	546.00
ENSG000000000457	339.75	316.50
ENSG000000000460	97.25	78.75
ENSG000000000938	0.75	0.00

Q5 (a). Create a scatter plot showing the mean of the treated samples against the mean of the control samples. Your plot should look something like the following.

Make a plot of control vs treated mean values for all genes

```
plot(meancounts)
```




Q5 (b). You could also use the ggplot2 package to make this figure producing the plot below. What geom\_?() function would you use for this plot?

```
library(ggplot2)
```

```
Warning: package 'ggplot2' was built under R version 4.5.2
```

```
ggplot(meancounts, aes(control.means, treated.means)) +  
  geom_point()
```



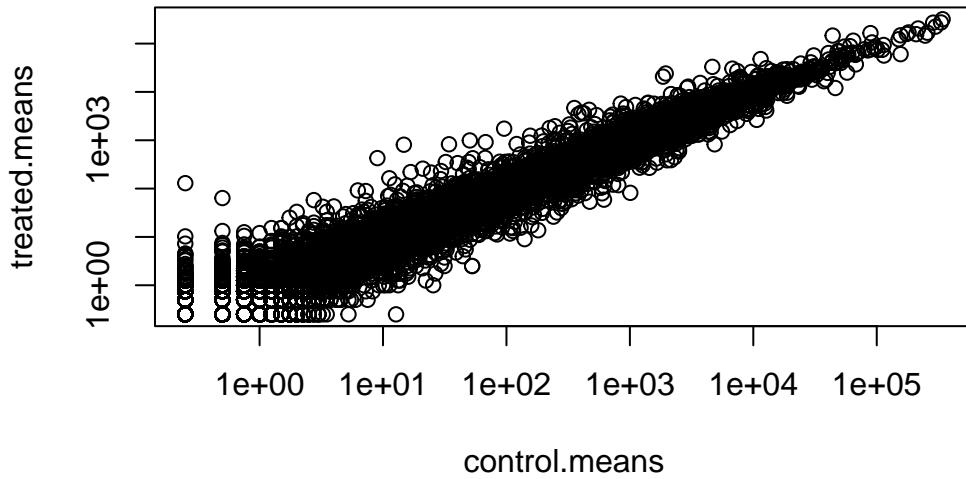
Q6. Try plotting both axes on a log scale. What is the argument to plot() that allows you to do this?

Make this a log plot

```
plot(meancounts, log="xy")
```

Warning in xy.coords(x, y, xlabel, ylabel, log): 15032 x values <= 0 omitted from logarithmic plot

Warning in xy.coords(x, y, xlabel, ylabel, log): 15281 y values <= 0 omitted from logarithmic plot



We often talk about metrics like “log2 fold-change”

```
# treated/control
log2(10/10)
```

```
[1] 0
```

```
log2(10/20)
```

```
[1] -1
```

```
log2(20/10)
```

```
[1] 1
```

```
log2(40/10)
```

```
[1] 2
```

```
log2(10/40)
```

```
[1] -2
```

Let's calculate the log2 fold change for our treated over control counts.

```
meancounts$log2fc <-  
log2( meancounts$treated.means / meancounts$control.means )
```

```
head(meancounts)
```

	control.means	treated.means	log2fc
ENSG000000000003	900.75	658.00	-0.45303916
ENSG000000000005	0.00	0.00	NaN
ENSG000000000419	520.50	546.00	0.06900279
ENSG000000000457	339.75	316.50	-0.10226805
ENSG000000000460	97.25	78.75	-0.30441833
ENSG000000000938	0.75	0.00	-Inf

```
zero.vals <- which(meancounts[ , 1:2]==0, arr.ind=TRUE)  
  
to.rm <- unique(zero.vals[,1])  
mycounts <- meancounts[-to.rm, ]  
head(mycounts)
```

	control.means	treated.means	log2fc
ENSG000000000003	900.75	658.00	-0.45303916
ENSG000000000419	520.50	546.00	0.06900279
ENSG000000000457	339.75	316.50	-0.10226805
ENSG000000000460	97.25	78.75	-0.30441833
ENSG000000000971	5219.00	6687.50	0.35769358
ENSG00000001036	2327.00	1785.75	-0.38194109

Q7. What is the purpose of the arr.ind argument in the which() function call above? Why would we then take the first column of the output and need to call the unique() function?

The argument arr.ind() in the which() function tells R to give the vector position of the TRUE values. The unique() function will return only 1 of that value instead of giving duplicates.

A common “rule of thumb” is a log2 fold change has a cutoff of +2 and -2 to call genes “Up regulated” or “Down regulated”.

Q8. Determine how many up regulated genes we have at the greater than 2 fc level.

Number of “up” genes

```
sum(meancounts$log2fc > +2, na.rm=TRUE)
```

```
[1] 1846
```

Q9. Determine how many down regulated genes we have at the greater than 2 fc level.

Number of “down” genes at -2 threshold

```
sum(meancounts$log2fc < -2, na.rm=TRUE)
```

```
[1] 2212
```

Q10. Do you trust these results? Why or why not?

I do not trust these results. This method of calculating the number of “up” and “down” genes fails to consider possible outliers/values that are not statistically significant by including these values in the results.

## DESeq2 analysis

Let’s do this analysis properly and keep our inner stats nerd happy - i.e. are the differences we see between drug and no drug significant given the replicate experiments.

```
library(DESeq2)
```

```
Warning: package 'matrixStats' was built under R version 4.5.2
```

For DESeq2 analysis, we need three things

- count values (`countData`)
- metadata tells us about the columns in `countData` (`colData`)
- design of the experiment (i.e. what do you want to compare)

Our first function from DESeq2 will setup the input required for analysis by storing all these three things together.

```
dds <- DESeqDataSetFromMatrix(countData=counts,
                               colData=metadata,
                               design=~dex)
```

converting counts to integer mode

Warning in DESeqDataSet(se, design = design, ignoreRank): some variables in  
design formula are characters, converting to factors

The main function in DESeq2 that runs the analysis is called `DESeq()`

```
dds <- DESeq(dds)
```

estimating size factors

estimating dispersions

gene-wise dispersion estimates

mean-dispersion relationship

final dispersion estimates

fitting model and testing

```
res <- results(dds)
head(res)
```

log2 fold change (MLE): dex treated vs control

Wald test p-value: dex treated vs control

DataFrame with 6 rows and 6 columns

	baseMean	log2FoldChange	lfcSE	stat	pvalue
	<numeric>	<numeric>	<numeric>	<numeric>	<numeric>
ENSG000000000003	747.194195	-0.3507030	0.168246	-2.084470	0.0371175
ENSG000000000005	0.000000		NA	NA	NA
ENSG000000000419	520.134160	0.2061078	0.101059	2.039475	0.0414026
ENSG000000000457	322.664844	0.0245269	0.145145	0.168982	0.8658106
ENSG000000000460	87.682625	-0.1471420	0.257007	-0.572521	0.5669691

```

ENSG000000000938    0.319167      -1.7322890  3.493601 -0.495846  0.6200029
                    padj
                    <numeric>
ENSG000000000003    0.163035
ENSG000000000005      NA
ENSG000000000419    0.176032
ENSG000000000457    0.961694
ENSG000000000460    0.815849
ENSG000000000938      NA

```


## Volcano plot

This is a common summary result figure from these types of experiments and plots the log2 fold change vs. the adjusted p-value.

```

plot(res$log2FoldChange, -log(res$padj))
abline(v=c(-2,2), col="red")
abline(h=-log(0.05), col="red")

```



## Save our results

```
write.csv(res, file="my_results.csv")
```

## Add gene annotation

To help make sense of our results and communicate them to others, we need to add some more annotations to our main `res` object.

We will use two bioconductor packages to first map IDs to different formats including the classic gene “symbol” gene name.

I will install these with the following commands: `BiocManager::install("AnnotationDbi")`  
`BiocManager::install("org.Hs.eg.db")`

```
library(AnnotationDbi)
library(org.Hs.eg.db)
```

Let's see what is in `org.Hs.eg.db` with the `columns()` function:

```
columns(org.Hs.eg.db)
```

```
[1] "ACNUM"      "ALIAS"       "ENSEMBL"      "ENSEMLPROT"   "ENSEMLTRANS"
[6] "ENTREZID"    "ENZYME"      "EVIDENCE"     "EVIDENCEALL"  "GENENAME"
[11] "GENETYPE"    "GO"          "GOALL"        "IPI"          "MAP"
[16] "OMIM"        "ONTOLOGY"    "ONTOLOGYALL" "PATH"         "PFAM"
[21] "PMID"        "PROSITE"     "REFSEQ"       "SYMBOL"       "UCSCKG"
[26] "UNIPROT"
```

We can translate or “map” IDs between any of these 26 databases using the `mapIds()` function.

```
res$symbol <- mapIds(keys = row.names(res), # our current IDs
                      keytype = "ENSEMBL",      # the format of our IDs
                      x = org.Hs.eg.db,        # where to get the mappings from
                      column = "SYMBOL")       # the format/DB to map to
```

```
'select()' returned 1:many mapping between keys and columns
```

```
head(res)
```

```
log2 fold change (MLE): dex treated vs control
Wald test p-value: dex treated vs control
DataFrame with 6 rows and 7 columns
  baseMean log2FoldChange    lfcSE     stat   pvalue
  <numeric>      <numeric> <numeric> <numeric> <numeric>
ENSG000000000003 747.194195 -0.3507030  0.168246 -2.084470 0.0371175
ENSG000000000005  0.000000      NA       NA       NA       NA
ENSG000000000419 520.134160  0.2061078  0.101059  2.039475 0.0414026
ENSG000000000457 322.664844  0.0245269  0.145145  0.168982 0.8658106
ENSG000000000460  87.682625 -0.1471420  0.257007 -0.572521 0.5669691
ENSG000000000938  0.319167 -1.7322890  3.493601 -0.495846 0.6200029
  padj      symbol
  <numeric> <character>
ENSG000000000003 0.163035    TSPAN6
ENSG000000000005  NA          TNMD
ENSG000000000419 0.176032    DPM1
ENSG000000000457 0.961694    SCYL3
ENSG000000000460 0.815849    FIRRM
ENSG000000000938 NA          FGR
```

```
Add the mappings for "GENENAME" and "ENTREZID" and store as res$genename and res$entrez
```

```
res$genename <- mapIds(keys = row.names(res),
  keytype = "ENSEMBL",
  x = org.Hs.eg.db,
  column = "GENENAME"
)
```

```
'select()' returned 1:many mapping between keys and columns
```

```
res$entrez <- mapIds(keys = row.names(res),
  keytype = "ENSEMBL",
  x = org.Hs.eg.db,
  column = "ENTREZID"
)
```

```
'select()' returned 1:many mapping between keys and columns
```

```
head(res)
```

```
log2 fold change (MLE): dex treated vs control
Wald test p-value: dex treated vs control
DataFrame with 6 rows and 9 columns
  baseMean log2FoldChange    lfcSE      stat     pvalue
  <numeric>      <numeric> <numeric> <numeric> <numeric>
ENSG000000000003 747.194195 -0.3507030  0.168246 -2.084470 0.0371175
ENSG000000000005  0.000000       NA        NA        NA        NA
ENSG000000000419 520.134160  0.2061078  0.101059  2.039475 0.0414026
ENSG000000000457 322.664844  0.0245269  0.145145  0.168982 0.8658106
ENSG000000000460 87.682625  -0.1471420  0.257007 -0.572521 0.5669691
ENSG000000000938 0.319167  -1.7322890  3.493601 -0.495846 0.6200029
  padj      symbol      genename      entrez
  <numeric> <character> <character> <character>
ENSG000000000003 0.163035   TSPAN6      tetraspanin 6      7105
ENSG000000000005  NA        TNMD       tenomodulin 64102
ENSG000000000419 0.176032   DPM1 dolichyl-phosphate m.. 8813
ENSG000000000457 0.961694   SCYL3 SCY1 like pseudokina.. 57147
ENSG000000000460 0.815849   FIRRM FIGNL1 interacting r.. 55732
ENSG000000000938  NA        FGR FGR proto-oncogene, .. 2268
```

## Pathway analysis

There are lots of bioconductor packages to do this type of analysis. For now, let's just try one called **gage** again we need to install this if we don't have it already.

```
library(gage)
library(gageData)
library(pathview)
```

To use **gage**, I need two things

- a named vector of fold-change values for our DEGs (our geneset of interest)
- a set of pathways or geneset to use for annotation

```
x <- c("barry"=5, "lisa"=10)
x
```

```

barry  lisa
      5     10

names(x) <- c("low", "high")
x

low high
5   10

foldchanges <- res$log2FoldChange
names(foldchanges) <- res$entrez
head(foldchanges)

7105       64102       8813       57147       55732       2268
-0.35070302        NA  0.20610777  0.02452695 -0.14714205 -1.73228897

data(kegg.sets.hs)

keggres = gage(foldchanges, gsets=kegg.sets.hs)

```

In our results object we have:

```

attributes(keggres)

$names
[1] "greater" "less"    "stats"

head(keggres$less, 5)

```

	p.geomean	stat.mean
hsa05332 Graft-versus-host disease	0.0004250461	-3.473346
hsa04940 Type I diabetes mellitus	0.0017820293	-3.002352
hsa05310 Asthma	0.0020045888	-3.009050
hsa04672 Intestinal immune network for IgA production	0.0060434515	-2.560547
hsa05330 Allograft rejection	0.0073678825	-2.501419
	p.val	q.val
hsa05332 Graft-versus-host disease	0.0004250461	0.09053483
hsa04940 Type I diabetes mellitus	0.0017820293	0.14232581
hsa05310 Asthma	0.0020045888	0.14232581

hsa04672 Intestinal immune network for IgA production	0.0060434515	0.31387180
hsa05330 Allograft rejection	0.0073678825	0.31387180
	set.size	exp1
hsa05332 Graft-versus-host disease	40	0.0004250461
hsa04940 Type I diabetes mellitus	42	0.0017820293
hsa05310 Asthma	29	0.0020045888
hsa04672 Intestinal immune network for IgA production	47	0.0060434515
hsa05330 Allograft rejection	36	0.0073678825

Let's look at one of these pathways (hsa05310 Asthma) with our genes colored up so we can see the overlap


```
pathview(pathway.id = "hsa05310", gene.data = foldchanges)
```

```
'select()' returned 1:1 mapping between keys and columns
```

Info: Working in directory C:/BIMM143/class12

Info: Writing image file hsa05310.pathview.png

Add this pathway figure to our lab report



## Save our main results

```
write.csv(res, file = "myresults_annotated.csv")
```